

Linice

A Linux Kernel Level Debugger

Version 2.0

www.linice.com

Author: Goran Devic

Contact email: author@linice.com

System Requirements

- Linux PC/x86 platform
- Minimum Pentium class CPU
- Linux kernel 2.4

Linice debugger has been developed on a number of Linux kernels, mainly on RedHat distributions. It was tested on RedHat 9.0 and SuSE 8.0. It should also work with no problems on some earlier versions of these distributions. Other distributions may work, and are not tested. You may have various successes with out-of-stock kernels or those earlier than 2.4. Linice currently does not work on kernels 2.6 (and 2.5).

Linice components

Linice consists of:

- 1) Linice kernel loadable module: kernel independent portion and kernel dependent portion. Build a generic object from within the “linice” directory. Kernel dependent loadable module is created by running a make within the “bin” directory.
- 2) Linsym: User mode app that loads / unloads Linice kernel module, generates and manipulates symbol files.
- 3) Xice: Initiates a session on the X-Server with a loaded Linice.

Compiling Linice from the source

The complete package can be compiled by running “make_bin” from the build directory. Individual components can be compiled by running “make” from within “linice”, “linsym” or “x” directories.

You will also need assembler NASM to successfully compile Linice.

When done compiling individual components, follow the instructions below to complete the kernel-dependent portion of the installation.

Compiling Linice from a binary package

- 1) Linice package comes with the pre-compiled kernel-independent object module “linice_kernel.o”. Before loading, it has to be linked with the iceface.c file which

contains your particular kernel interfaces. You will have to run “make”; but first, you may need to adjust some configuration defines:

Modify ‘Makefile’ to suit your environment:

Default:	TARGET=
Configuration:	Single CPU, non- APIC

By default, Linice builds for a single CPU, non-APIC target machine.
If your machine has IO_APIC, you will need to add `-DIO_APIC` define.
If your machine is SMP, you will need to add `-DSMP` define.

Default:	#PCIHDR = PCIHDR
Configuration:	Include PCI header information file

Delete “#” in order to include the “pcihdr.h” file. This file contains a database of PCI devices. It is used by the PCI command to decode devices on a PCI bus. Inclusion of this file will enlarge the module by about 300K.

2) Run ‘make’ in order to generate the debugger module tailored to your running kernel.

3) Customize initialization file “linice.dat”

Specify the memory to be reserved for symbols, history buffer or display; change the init string or macros, or set the different keyboard layout to customize it to your particular locale. The default values are a good starting point and should work.

These keyboard layouts are supported:

us	United States (default)
uk	UK
finnish, finnish-latin1	Finland
de, de-latin1	Germany
fr, fr-latin1	France
dk, dk-latin1	Denmark
dvorak	[Dvorak keyboard]
sg, sg-latin1	Switzerland (German)
sf, sf-latin1	Switzerland (French)
be	Belgium
po	Portugal
it	Italy
sw	Sweden
hu	Hungary
jp106	Japan
pl	Poland
hr-cp852, hr-latin2	Croatia / Hrvatska
cz-qwerty, cz-qwertz	Czech Republic

It is necessary to let Linice know of a proper keyboard layout since it does not use Linux code for keyboard handling.

- 4) Use linsym to load and unload the Linice module:

Load and install Linice kernel module: “linsym -i”

Unload Linice: “linsym -x”

You have to use linsym to load (and unload) Linice since Linice needs some initial symbols that linsym provides. Also, linsym reads the initialization file “linice.dat” and sends it as a part of the initial packet.

Linsym searches for the kernel symbols in the /boot/System.map. That symbol map file has to match the running kernel! This is very important since the linsym reads some kernel addresses from it that helps it to locate code to place various hooks (such as a keyboard hook or a pointer to the kernel module list.)

- 5) Use linsym to translate symbols

Translate debug information from a program or a module: “linsym -t <binary>”

This will create a symbol file <binary>.SYM that can be loaded.

- 6) Use linsym to load symbol file

Load symbol file into Linice: “linsym -s <symbols.sym>”

- 7) You may break into the Linice at any time by pressing the keyboard hotkey: “Ctrl+Q”

Running in the X-Window

To have the Linice pops up on the X-Window, you have to have it loaded using the standard command “linsym -i”. If you are in the X environment already, it might look as the system froze since the Linice actually popped up using the VGA frame buffer. Simply hit F5 (go) and you will get back the terminal control. Then simply run the utility “xice”, and the Linice will appear on the top of your X window.

Compiling the debuggee

Notes on compiling your code to be debugged:

Program or module has to be compiled with the “-gstabs+” switch in order to generate symbolic information suitable for Linice translation.

To get the proper visibility into local symbols, you should **not** use the switch: “-fomit-frame-pointer” since with that switch the local variable information will not be accessible.

It is recommended to disable code optimization.

To get the visibility into non-static, un-initialized global symbols, you need to use the linker switch: “-dp” or “-d” or “-dc” to force the assignments of space for “common symbols”.

Preparing the Linux kernel for source debugging

Although you can run Linice on top of the unmodified kernel, if you need to do source level debug on the kernel code, these steps should give you a general guidance on how to prepare the kernel code and symbols. You will need to rebuild the kernel in order to generate necessary stabs info, which will be used by the symbol translator. Edit the Linux kernel Makefile and insert the code in red/underline:

We need stabs debugging information to be built with the Linux kernel. Add “-gstabs+” to the CFLAGS and add “-gstabs” to the AFLAGS:

```
CFLAGS := -gstabs+ $(CPPFLAGS) -Wall -Wstrict-prototypes -Wno-trigraphs
-O2 -fno-strict-aliasing -fno-common -Wno-unused
AFLAGS := -gstabs -D__ASSEMBLY__ $(CPPFLAGS)
```

In order to be able to access local variables, remove the option *-fomit-frame-pointer*.

Next, we will make the default build generate an intermediate kernel version with all the stabs information, then have them stripped for the version that we will be loading:

```
$(LD_VMLINUX) $(LD_VMLINUX_KALLSYMS) -o vmlinux.debug
$(STRIP) -S -o vmlinux vmlinux.debug
$(NM) vmlinux | grep -v '\\(compiled\\)\\|\\(\\.o\\$\\)\\|\\( [aUw]
\\)\\|\\(\\.\\.ng\\$\\)\\|\\(LASH[RL]DI\\)' | sort > System.map
```

This modified Linux Makefile will now generate the file vmlinux.debug, which will be used to generate the Linice symbol file. It will also generate a regular kernel code that will be installed. Follow the standard kernel build procedure:

```
# make mrproper
# make xconfig or make menuconfig or make config
# make dep
# make clean
# make bzImage
```

vmlinux.debug and vmlinux images are now built. vmlinux is a stripped version of the vmlinux.debug.

```
# make modules
# make modules_install
# make install
```

Translate vmlinux.debug kernel symbols into Linice symbol file:

```
# linsym -t vmlinux.debug
```

Add new kernel to the boot loader and reboot using it, so the new symbols will match the running kernel when you load them into Linice.

A note about the serial connection

When you use a serial port and the VT100 terminal is connected to the other side, the local keyboard is still being used as an active input device. Although that behavior appears odd, it helps to keep control when the connection parameters are not quite right. This behavior may be changed in the future.

Some Toshiba notebooks have a hidden serial port connector under the keyboard. It uses a non-standard IO port of 0x1E0. Linice supports it and the port is enumerated as COM5.

Bugs

Please send all bug reports to bugs@linice.com.

Include Linice version, the version of the kernel that you are running it on (or trying to run it on), and any other information that would help reproduce the problem. You may include file “Version.txt” that is part of your Linice distribution; it contains the Linice version and the build date.

Debugging and developing Linice

If you are inclined to help out and work on Linice, please send me a note and I will try to help you get going as best as I can.

Since all the sources are enclosed, you are of course free to tinker with it. Please send me all modifications that you deem useful for others, so I can add them to the “official” package for everyone’s benefit. Of course, you will be credited as well ☺. You can send me modified source file(s) or just diffs.

Features and wish list

If you would like to see a feature implemented or have other suggestion for improvement, please email to features@linice.com.

USB Keyboard not supported

This limitation stems from the fact that Linice is handling keyboard at the low-level and currently only knows how to handle a legacy PS/2 interface.

If you simply cannot plug in the PS/2 keyboard (which would solve this problem), and you really have to use a USB keyboard, you will need to disable Linux support for all USB devices. This will allow the system BIOS to handle the USB keyboard so it will appear to the software as if you have a legacy PS/2 keyboard. If you enable any USB support in the Linux kernel, it will turn off SBIOS handling of the USB keyboard, and Linice will not be able to handle it.

New commands

These commands are added and don’t exist on the SoftIce©:

ASCII

Prints an ASCII character table

XWIN

Redirect console to a DGA frame buffer (X-Window)

Not implemented commands

These commands are not (yet) ported from the SoftIce©:

A (Assemble code)
ADDR (Display/change address contexts)
BH (Breakpoint history)
BPINT (Breakpoint on interrupt)
DATA (Change data window)
DEVICE (Display info about a device)
DEX (Display/assign window data expressions)
GENINT (Generate an interrupt)
PAGEIN (Load a page)
PRN (Set printer output)
QUERY (Display a process virtual address space map)
SHOW (display from backtrace buffer)
SS (Search source module for string)
STACK (Display call stack)
THREAD (Show thread information)
TRACE (Enter back trace simulation mode)
WS (Toggle call stack window)
XG (Trace simulation)
XP (Step in trace simulation)
XRSET (Reset trace history buffer)
XT (Step in trace simulation)
XFRAME (Display active exception frames)

Implemented commands

SETTING BREAK POINTS

BPM - Breakpoint on memory access
BPMB - Breakpoint on memory access, byte size
BPMW - Breakpoint on memory access, word size
BPMD - Breakpoint on memory access, double word size
BPIO - Breakpoint on I/O port access
BPX - Breakpoint on execution
BSTAT - Breakpoint Statistics

MANIPULATING BREAK POINTS

BPE - Edit breakpoint
BPT - Use breakpoint as a template
BL - List current breakpoints

BC - Clear breakpoint
 BD - Disable breakpoint
 BE - Enable breakpoint
DISPLAY/CHANGE MEMORY
 R - Display/change register contents
 U - Un-assembles instructions
 D - Display memory
 DB - Display memory, byte size
 DW - Display memory, word size
 DD - Display memory, double word size
 E - Edit memory
 EB - Edit memory, byte size
 EW - Edit memory, word size
 ED - Edit memory, double word size
 PEEK - Read from physical address
 PEEKB - Read from physical address a byte
 PEEKW - Read from physical address a word
 PEEKD - Read from physical address a dword
 POKE - Write to physical address
 POKEB - Write to physical address a byte
 POKEW - Write to physical address a word
 POKED - Write to physical address a dword
 H - Help on the specified function
 HELP - Help on the specified function
 ? - Evaluate expression
 VER - Linice version
 WATCH - Add watch variable
 FORMAT - Change format of data window
DISPLAY SYSTEM INFORMATION
 GDT - Display global descriptor table
 LDT - Display local descriptor table
 IDT - Display interrupt descriptor Table
 TSS - Display task state segment
 CPU - Display cpu register information
 PCI - Display PCI device information
 MODULE - Display kernel module list
 PAGE - Display page table information
 PHYS - Display all virtual addresses for physical address
 PROC - Display process information
 WHAT - Identify the type of an expression
I/O PORT COMMANDS
 I - Input data from I/O port
 IB - Input data from I/O port, byte size
 IW - Input data from I/O port, word size
 ID - Input data from I/O port, double word size
 O - Output data to I/O port
 OB - Output data to I/O port, byte size
 OW - Output data to I/O port, word size
 OD - Output data to I/O port, double word size
FLOW CONTROL COMMANDS
 X - Return to host debugger or program
 G - Go to address
 T - Single step one instruction
 P - Step skipping calls, Int, etc.
 HERE - Go to current cursor line
 HALT - System APM Off
 HBOOT - System boot (total reset)

MODE CONTROL

11HERE - Direct INT1 to LinICE, globally or kernel only
I3HERE - Direct INT3 to LinICE, globally or kernel only
ZAP - Zap embedded INT1 or INT3
FAULTS - Enable/disable LinICE fault trapping
SET - Change an internal system variable
VAR - Change a user variable

CUSTOMIZATION COMMANDS

PAUSE - Controls display scroll mode
ALTKEY - Set key sequence to invoke window
FKEY - Display/set function keys
CODE - Display instruction bytes in code window
COLOR - Display/set screen colors
TABS - Set/display tab settings
LINES - Set/display number of lines on screen
WIDTH - Set/display number of columns on screen
MACRO - Define a named macro command

UTILITY COMMANDS

S - Search for data
F - Fill memory with data
M - Move data
C - Compare two data blocks
ASCII - Prints an ASCII character table

LINE EDITOR KEY USAGE

up - Recall previous command line
down - Recall next command line
right - Move cursor right
left - Move cursor left
BKSP - Back over last character
HOME - Start of line
END - End of line
INS - Toggle insert mode
DEL - Delete character
ESC - Cancel current command

WINDOW COMMANDS

WC - Toggle code window
WD - Toggle data window
WL - Toggle locals window
WR - Toggle register window
WW - Toggle watch window
EC - Enter/exit code window
.
- Locate current instruction

WINDOW CONTROL

VGA - Switch to a VGA text display
MDA - Switch to a MDA (Monochrome) text display
XWIN - Redirect console to a DGA frame buffer
SERIAL - Redirect console to a serial terminal
CLS - Clear window
RS - Restore program screen
ALTSCR - Change to alternate display
FLASH - Restore screen during P and T

SYMBOL/SOURCE COMMANDS

SYM - Display symbols
EXP - Display exported symbols from a kernel or a module
SRC - Toggle between source, mixed & code
TABLE - Select/remove symbol table
FILE - Change/display current source file

TYPES - List all types, or display type definition
LOCALS - Display locals currently in scope
SPECIAL OPERATORS
. - Preceding a decimal number specifies a line number
@ - Preceding an address specifies indirection

Linsym – Symbol loader and translator

The following is a list of supported arguments (or options) to the linsym utility.

Option: **--install**

Short option: **-i**

Installs Linice debugger module and breaks.

Example: `# linsym -i`

Option: **--map <System.map>**

Short option: **-m <System.map>**

In order to successfully load Linice, Linsym needs a current System.map file, which it will try to find at certain default locations (/boot/System.map, /boot/System.map-<kernel name>). If you have custom-compiled your kernel, the current system map may be at a different location, or be named differently. Use this option when loading Linice to specify the correct path and name of that file, if necessary.

Example: `# linsym -m /boot/System.map-test -i`

Option: **--uninstall**

Short option: **-x**

Uninstalls Linice debugger module.

Example: `# linsym -x`

Option: **--translate <program>**

Short option: **-t <program>**

Translates stabs symbols from your module, kernel or executable program and creates a separate symbol file. This symbol file contains all the available debug information that was compiled and linked with the target program: global and local symbols, source code, type definition and other pertinent debugging information.

Example: `# linsym -t module.o`

Option: **--output <alt_name.sym>**

Short option: **-o <alt_name.sym>**

Specifies alternate file name for the symbol file generated by the translation (option “-t”).

Example: `# linsym -t module.o -o symbol.sym`

Option: **--path <orig-path>:<new-path>**

Short option: **-p <orig-path>:<new-path>**

Specifies path substitution for the source code. This option is useful when you are building a symbol file from another computer, and the source code that you would include resides on a different directory path. This option lets you substitute a path prefix. Note that the two paths are separated by a colon: the first path is the path that will be matched against all original absolute paths in the stabs debugging section (specifying the path to the source), and the second path is the path to be used instead.

Example: # `linsym -t module -p /usr/src/mod:/mnt/usr/src/mod`

Option: --sym <symbol.sym>

Short option: `-s <symbol.sym>`

Loads one or more symbol files into the running Linice. You can load multiple symbol files by separating them with a colon “:”. The names of symbol files usually end with the “.sym”, which has to be specified as part of the file name. This is in contrast to the “unload” command, which specifies the module or program name.

Example: # `linsym -s module.sym`

Option: --unload <symbol.sym>

Short option: `-u <symbol>`

Unloads one or more symbol files from the running Linice. You can unload multiple symbol files by separating them with a colon “:”. You can also unload symbol files from within the Linice using the command “table”. The specified name is the name of the symbol table as listed in the Linice using the “table” command. (It is not the file name which should be used when loading a symbol file.)

Example: # `linsym -u module`

Option: --logfile [<filename>][,append]

Short option: `-l [<filename>][,append]`

Saves the content of the Linice history buffer (the command line window) into a file. You can optionally provide a file name; if you don’t, the default file name “linice.log” will be used. If the file already exists, it will be truncated, unless you specified “,append” also, which will preserve the original content and append a new one.

Example: # `linsym -l out.log,append`

Option: --verbose {0-3}

Short option: `-v {0-3}`

Specifies verbose level. Default is 0, which is silent. If Linice refuses to load or you encounter some other error, set the verbose level to 3 and see if the message dump helps. If not, send me an email with the dump attached. The most common problem is incorrect system.map file (if you have compiled a custom kernel).

Linice configuration file: linice.dat

Configuration file “linice.dat”, which has to reside either in the current directory, or in the /etc directory, contains initial parameters and switches for the loading Linice debugger. For most users, it should work just fine without modifications.

This text file is fully compatible with the SoftIce version. This section describes the parameter values that are used by Linice (the rest of the tokens are ignored; they may be used in the future versions of Linice)

lowercase = [on | off]

Specifies the initial disassembly character case.

sym = <buffer in Kb>

Linice will reserve this buffer for all symbol files that will be loaded. Specify it large enough since it cannot be modified without reloading the debugger.

hst = <buffer in Kb>

Specify the size of the history buffer (which is also known as a command line buffer.) If you intend to save this buffer to a file, you may want to give it a larger size.

macros = <number>

Number of keyboard macros that are going to be allocated.

drawsize = <buffer in Kb>

If you will use Linice on the X-Window display, this buffer will be allocated to store the background frame buffer at the location that Linice overwrites.

init = <init commands;>

This line specifies a set of Linice commands that will be executed immediately upon Linice load.

F1...F12 Function keys assignment

SF1...SF12 Shift + Function keys assignment

AF1...AF12 Alt + Function keys assignment

CF1...CF12 Ctrl + Function keys assignment

Each key in these combinations may be assigned a command or a set of commands.

layout = [country-code]

Specify the keyboard layout as a country code, which is described earlier. The default layout is US.

Advanced Topics – Debugger Extensions

Linice debugger supports custom plug-ins, so-called “dot-commands” (since they are typed after a dot/period on the command line).

Command: ? List all standard Linice commands
Command: .? List all registered custom dot commands
Command: .<cmd> Execute a registered dot command

You can write a debugger extension that implements one or more dot-commands as a kernel loadable module. There are 2 functions provided by Linice which are used to register and unregister extension interface.

Please refer to the header file `LiniceExt.h` for the details of the interface. There is also a sample module to illustrate the function and capability of the interface. You can compile it and run it, or use it as a starting point for your own extensions.

int LiniceRegisterExtension(TLINICEEXT *pExt);

Use this function to register a debugger extension interface. It returns one of the error codes specified in the `LiniceExt.h` header file. `TLINICEEXT` is a structure that describes the interface and contains all the callback function pointers.

The caller is responsible to initialize these portions of the structure `TLINICEEXT`:

version

Set it to the macro `LINICEEXTVERSION`. This field is mandatory. Linice is using it to verify the correct and supported interface version.

size

Set it to the macro `LINICEEXTSIZE`. This field is also mandatory.

pDotName

Set it to the ASCIIZ name of the dot command that you implement using this interface structure. This field is mandatory. Your code may support a number of extensions, each being represented by one `TLINICEEXT` structure.

pDotDescription

This field is optional. If not used, set it to `NULL`; otherwise, set it to the ASCIIZ string with the description of the command. This string will be used when a user lists all dot commands using the command “.?”. The description helps by telling what the command does, and serves no other purpose.

```
int ( *Command ) ( char *pCommand ) ;
```

Specify the address of your function which will be called when a dot command is invoked. This can be NULL, which would make sense only if you don't intent to provide a command handler, but only a function token handler (described below). Note that you still need to provide *pDotName* string even if you don't have a handler.

When this function gets called, argument *pCommand* points to the arguments of the command, or the rest of the line.

Example: If you registered a command "dump", and the user typed ".dump 1 2 3" on the command line, *pCommand* would point to " 1 2 3" (also note the space before "1".)

```
void (*Notify)(int Notification);
```

Specify the address of your function which will be called on various debugger system events. This can optionally be NULL if you don't care about the events.

The debugger events ("Notification") are:

PEXT_NOTIFY_ENTER - Linice got control. Break into debuggee.

PEXT_NOTIFY_LEAVE - Linice released control. Debuggee continues to run.

```
int (*QueryToken)(int *pResult, char *pToken, int len);
```

Specify the address of your function which will be called when an unknown token is encountered within an expression. Your extension may be able to help parsing expressions and provide values to add-on functions or tokens.

These are the parameters sent by Linice: *pToken* points to the start of the expression token, *len* specifies the suggested length of the token, and *pResult* points to a variable of the type 'int' where you should store the result.

After you examine a token (probably using a function 'strncmp()' or similar), if you detect that you don't handle it, simply return 0.

If you do handle the token, store the final value in the **pResult* and return the number of characters to advance past the token size. This may or may not be the same value as parameter *len*.

If the token is a function, it is possible to recursively call the expression evaluator for the function arguments in parenthesis. The sample extension module shows how to do it.

After the extension is successfully registered, Linice fills in the pointers to some of its utility functions that your module can call:

```
TLINICEREGS *(*GetRegs)(void);
```

Returns the address of the internal Linice structure that holds the CPU registers of the program being debugged. This address will not change for the duration of Linice session. You can read and write CPU registers when the debugger is active.

```
int (*Eval)(int *pValue, char *pExpr, char **ppNext);
```

This function evaluates a string expression into a number. One possible use is to resolve an expression that a user might have typed as arguments to your dot command.

Set *pValue* to where the result should be stored. Set *pExpr* to the expression string, and set *ppNext* to the char* variable to receive the end of the evaluated string. (*ppNext* is optional: if you don't care where the expression ended, set it to NULL. If you need to parse multiple expressions that are given one after the other, then you may want to know where the previous expression ended.)

On success, the function returns a non-zero value. On failure, the function returns 0, *pValue* is not modified, and *ppNext* points to the character which caused the error.

```
int (*Disasm)(char *pBuffer, int sel, int offset);
```

Disassemble a line of x86-code into your buffer from the given address. *pBuffer* needs to be at least 80 characters long to store disassembled instruction. *sel* and *offset* define the target address. If you specify 0 for *sel*, kernel CS will be used instead.

```
int (*dprint)(char *format, ...);
```

This function provides a way to print out any message with variable number of arguments into the Linice command buffer. It works similar to the standard “C” function `printf()`, except that you should not use special character for new line “\n”.

```
int (*Execute)(char *pCommand);
```

This function executes any command that you might also be able to type on the command line.

```
int (*Getch)(int fPolled);
```

Reads and returns a character from the input stream. Use this function for the interactive option menus within the extension handler. Set *fPolled* to TRUE to have the function wait until a key becomes available.

```
int (*MemVerify)(int sel, int offset, int size);
```

This function verifies that a range of memory addresses is present and accessible. When it returns a nonzero value, you can access the memory range using your pointers. *sel* and *offset* specify the start address, and *len* is the size of the memory range in bytes. If you specify 0 for *sel*, kernel DS will be used instead.

`void LiniceUnregisterExtension(TLINICEEXT *pExt);`

Use this function to unregister previously registered interface. Be sure to call this function before unloading your extension module, so the Linice stops calling it via the registered structure. Failure to do so will result in a crash.

You have to unload all of your extension modules before you unload Linice, since the modules are linked to Linice by the means of those two exported functions. Run the Linux command “`lsmod`” to see the module dependency. Alternatively, list the registered extensions with the Linice command “`.?`”.

Be sure not to call any of the callbacks from the extension interface structure if the registration failed.